

Containers – Introduction

RSE L&D 27/02/2024

Dr Franz Lang

Overview

- Big Picture
- VM vs Container
- Architecture
- Workflow
- Docker vs Singularity/Apptainer
- Apptainer definition file
- Examples
- Registry
- Container Orchestration

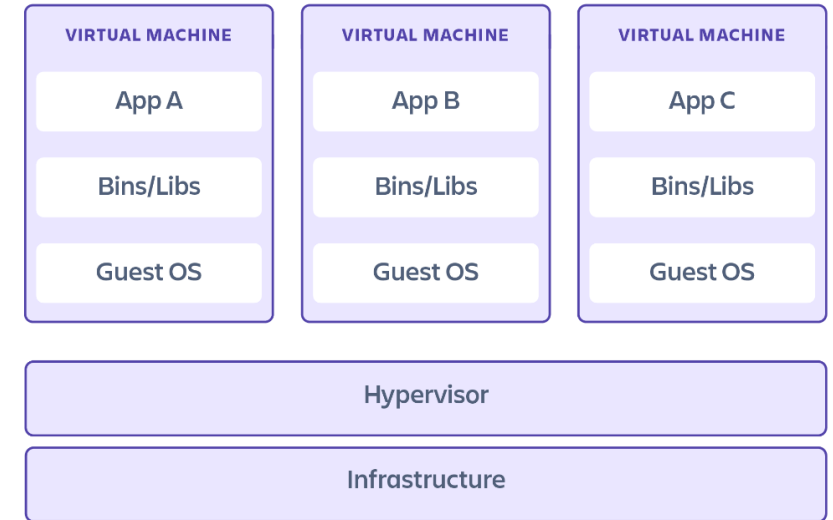
Big Picture

- Interoperability: OS-agnostic (almost)
 - Windows containers only on Windows
 - Linux containers on Windows via WSL
 - kernel function (in-)compatibilities
- Reproducible: standalone package for each app/service
 - contains all dependencies required
 - lightweight: can optimise content to be minimal
- Maintainable: text-file build recipe
- Deployable: hosting via a registry
 - dockerhub, harbour
- Orchestration: usable with clusters
 - kubernetes, docker swarm, podman...
- Security:
 - isolation of applications
 - can apply user privileges and security settings

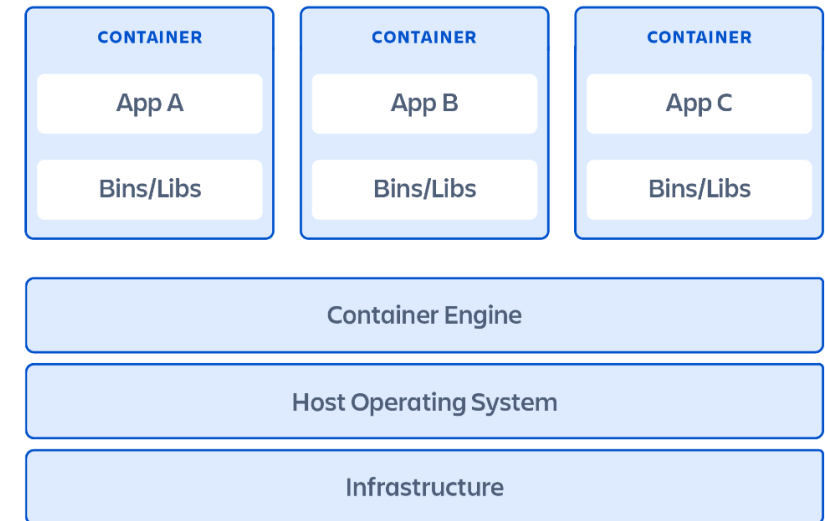
VM vs Container

- VM
 - abstraction of physical hardware: 1 server to many servers
 - hypervisors allows multiple VMs on single machine
 - each VM contains full OS, binaries, libraries, application
 - typically larger and slower to boot
- Container
 - abstraction at application layer
 - multiple containers can run on same machine as isolated processes
 - share host OS kernel
 - typically smaller and fast to launch
 - made up of layers

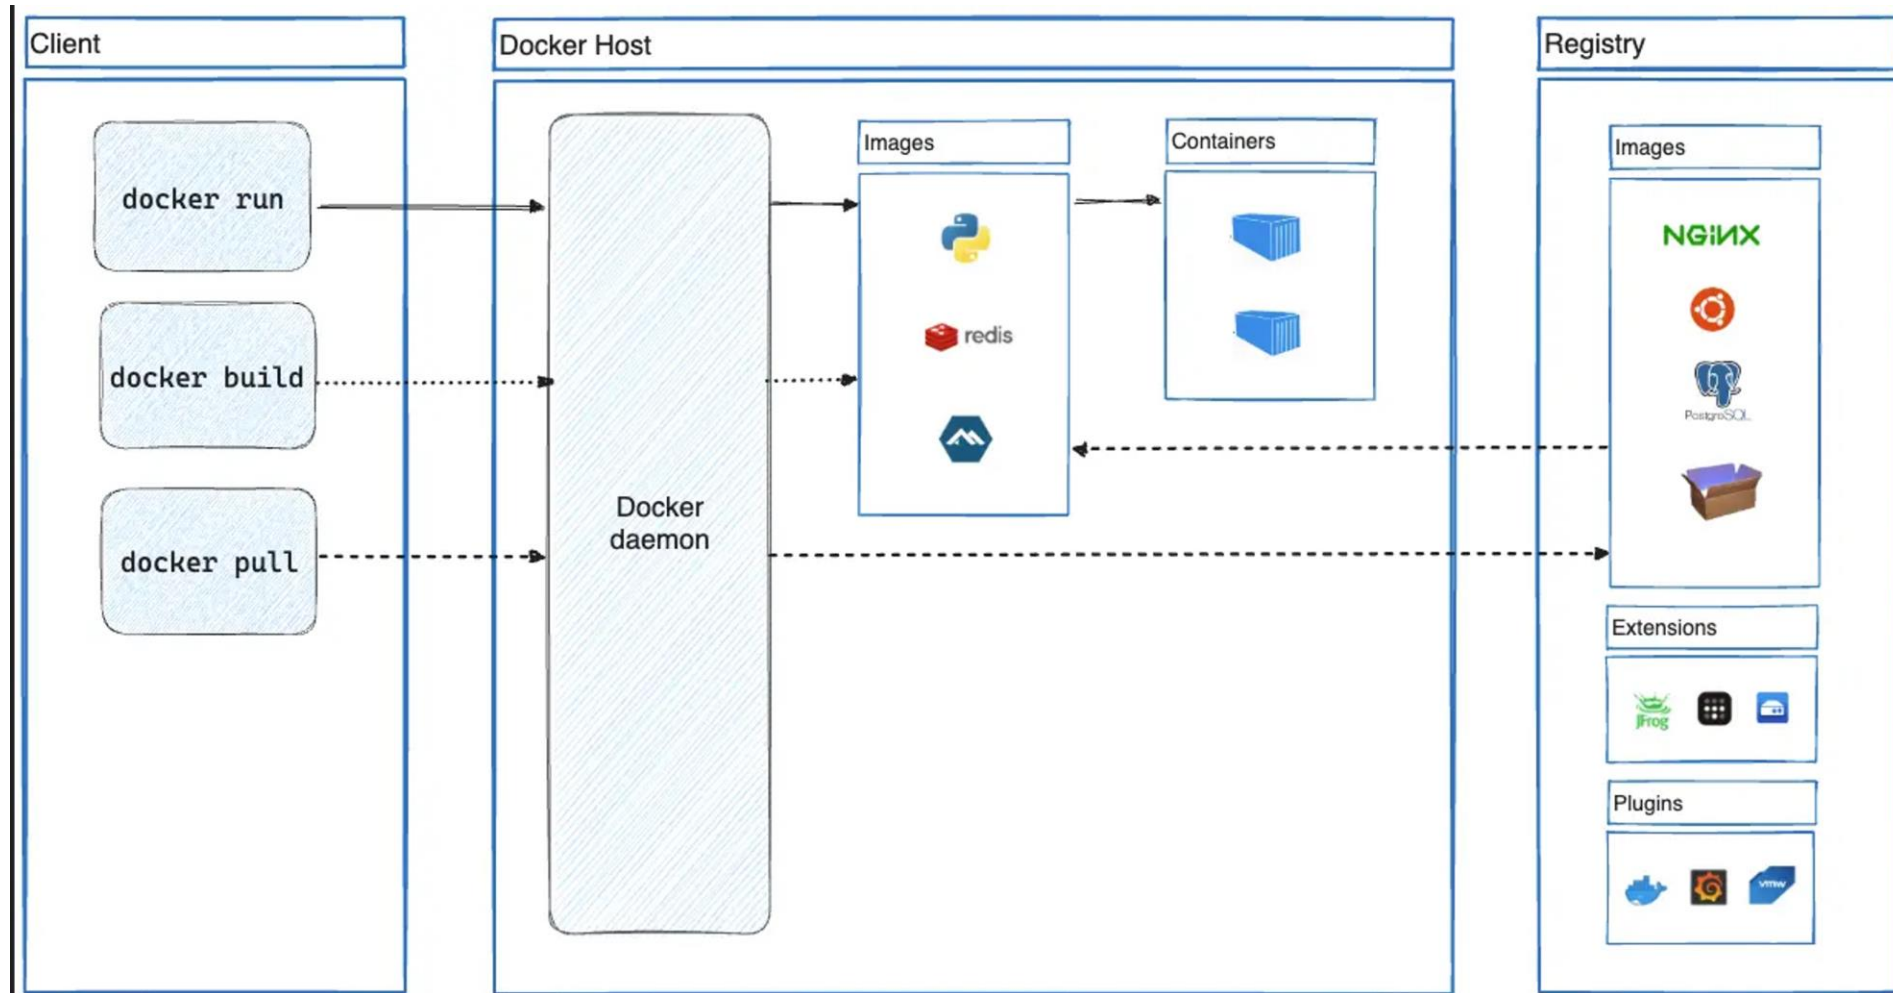
Virtual machines



Containers



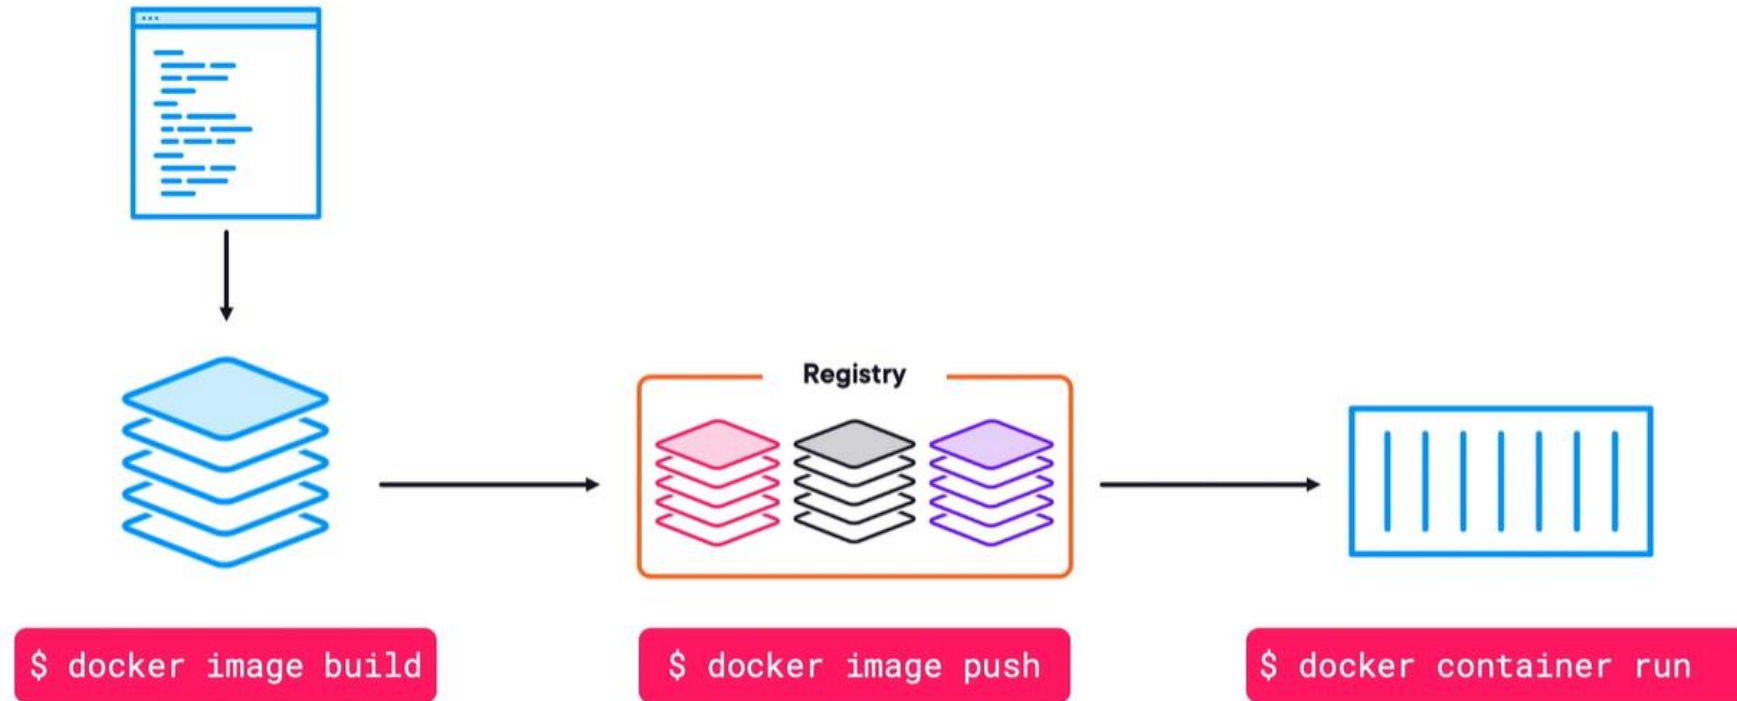
Architecture



Graphic from <https://docs.docker.com/get-started/overview/>

Workflow

- create starting build recipe
- create image in sandbox mode
- develop/debug image in sandbox mode
- update build recipe to final version
- create final image
- push image to registry
- test deployment of image via the registry



Docker vs Singularity/Apptainer

Docker

- + established orchestration options
 - + docker-compose, Kubernetes, podman, docker swarm
- + existing registry: Dockerhub
- + Windows containers possible
- requires root access (hard to use on HPC)
- Windows: Linux containers via WSL only
- commercial

Singularity/Apptainer

- + open source
- + no root required (usable on HPC)
- + can use all Docker containers
- Windows: Linux containers via WSL only
- registry creation
- no easy orchestration

Singularity project split in 2021 into two forks: Singularity CE (<https://sylabs.io/>) and Apptainer (<https://apptainer.org/>). The original Singularity project is no longer being developed.

Apptainer definition file - Header

- definition (or `def`) file divided into two parts: Header and Sections
- blueprint for building a container
- mandatory: `Bootstrap` and `From`
- Bootstrap
 - agent used for creating base OS
 - depends on the container/OS you want to use
- From
 - determines OS to use
- additional options for multi-stage builds or specific URLs

```
Bootstrap: docker
From: ubuntu:{{ VERSION }}
Stage: build
```


Apptainer definition file - Sections

- different sections prefaced by `%` symbol to add different content/commands/features during build process
- none mandatory
- possible section names:
 - arguments, setup, **files**, **app***, **post**, test, **environment**, startscript, **runscript**, labels, **help**
- order of sections within definition file irrelevant
- order of execution is well-defined: order of section names above

```
Bootstrap: docker
From: ubuntu:{{ VERSION }}
Stage: build

%arguments
VERSION=22.04

%setup
touch /file1
touch ${APPTAINER_ROOTFS}/file2

%files
/file1
/file1 /opt

%environment
export LISTEN_PORT=54321
export LC_ALL=C

%post
apt-get update && apt-get install -y netcat
NOW=`date`
echo "export NOW=\"${NOW}\"" >> $APPTAINER_ENVIRONMENT

%runscript
echo "Container was created $NOW"
echo "Arguments received: $*"
exec echo "$@"

%startscript
nc -lp $LISTEN_PORT

%test
grep -q NAME="Ubuntu" /etc/os-release
if [ $? -eq 0 ]; then
    echo "Container base is Ubuntu as expected."
else
    echo "Container base is not Ubuntu."
    exit 1
fi

%labels
Author alice
Version v0.0.1

%help
This is a demo container used to illustrate a def file that uses all
supported sections.
```

Examples

- login into IDAaaS and create a workspace
 - small CPU + RAM sufficient
 - e.g. Muon (EMU, MUSR, HIFI,...) or Excitations Powder (ALF, LET, MAPS, ...)
- open a terminal and git clone repo for session
 - git clone https://github.com/franzlang/container_workshop.git
- slides and examples (with solutions) contained in repo

Example 1 – hello world

- make sure you can run Apptainer by downloading and running Docker's `hello-world` container
 - for example via
apptainer run docker://hello-world
- learn to:
 - specify a OS
 - build your own minimal image
 - make your image perform a basic task
 - run your minimal image in a container
- %help
 - metadata for helpful information about the container
- %runscript
 - commands executed when container is run
 - \$*: options passed to container at runtime
 - \$@: passing the options to a command via a quoted array

```
$ apptainer build alpine.sif docker://alpine
```

```
$ apptainer build lolcow.sif lolcow.def
```

```
apptainer run [run options...] <container> [args...]
```

%help

This is a demo container used to illustrate a def file that uses all supported sections.

%runscript

```
echo "Container was created $NOW"  
echo "Arguments received: $*"  
exec echo "$@"
```

Example 2 – basic shell commands

- learn to:
 - use the sandbox feature to build/develop/debug a container interactively
 - open a shell inside a container to run commands inside the container
 - use %post to run shell commands when building the image
 - use %environment to specify environment variable(s)
- %post
 - commands run during image creation
- %environment
 - setting environment variables within the container

```
apptainer shell [shell options...] <container>
```

%post

```
apt-get update && apt-get install -y netcat
```

%environment

```
export LISTEN_PORT=54321  
export LC_ALL=C
```

Example 3 – files

- learn to:
 - use the %files section to add a file from the host system to the container image
- %files
 - copying files from a stage during the build to another stage
 - e.g.: compile stage creates binary, which is copied to final stage
 - copying files from host machine into the image

```
%files [from <stage>]  
    <source> [<destination>]  
    ...
```

```
%files  
    /file1  
    /file1 /opt
```

Example 4 – multi-stage build

- learn to:
 - use multiple build stages using the `Stage` part of the Header
 - use files from one build stage, e.g. compiled binaries, in another build stage
- advantage: can deploy compiled binaries without deploying build dependencies

Example 5 – multi-image build

- learn to:
 - use a custom-built image as the base image for a follow-on image
 - use a local image in the build process via the `From` part of the Header
- advantage: can separate complicated images into parts
 - e.g.: parts that rarely change (compiler, simulation engine, ...) and parts that change regularly

Preferred bootstrap agents 🔗

- [docker](#) (images hosted on Docker Hub)
- [oras](#) (images from supported OCI registries)
- [localimage](#) (images saved on your machine)
- [scratch](#) (a flexible option for building a container from scratch)

Other bootstrap agents

- [library](#) (images hosted on Library API Registries)
- [shub](#) (images hosted on Singularity Hub)
- [yum](#) (yum-based systems such as CentOS and Scientific Linux)

Example 6 – multi-app image

- learn to:
 - allow multiple applications to be callable from the same container
 - forward input to the application inside the container
- advantage: can run multiple binaries/apps from the same container

```
%apprun foo
    exec echo "RUNNING FOO"

%appstart foo
    exec echo "STARTING FOO"

%applabels foo
    BESTAPP FOO

%appinstall foo
    touch foo.exec

%appenv foo
    SOFTWARE=foo
    export SOFTWARE

%apphelp foo
    This is the help for foo.
```

```
% apptainer run --app foo my_container.sif
RUNNING FOO
```


Example 7 – mounting volume

- learn to:
 - mount a local volume of the host machine within the container
- advantage: have a shared filesystem between the container and the host

```
$ ls /data
bar  foo

$ aptainer exec --bind /data:/mnt my_container.sif ls /mnt
bar  foo
```

Example 8 – GUI/X11

- example to illustrate how a container can launch a graphical element

Example 9 – Jupyter

- learn to:
 - build a minimal python container to run a Jupyter notebook server
 - run a Jupyter notebook server in the container
 - access the server via the host's browser
 - run a notebook on the host system using the containerised notebook server

Real-life examples

- included a number of examples used to deploy applications on IDAaaS
- containers will take too long to build
 - or will not build if I had to sanitise the definition file
- look for the useful Apptainer/container feature that the example tries to illustrate

Registry

- upload images to a cloud/server to allow easier deployment
- STFC cloud has a Harbor registry
 - <https://harbor.stfc.ac.uk/>
 - login via IRIS IAM and fedID+password
- IDAaaS has dedicated repository
 - can gain push access to specific project
 - allows easy deployment on IDAaaS
- can hook registry into github actions for easy CI/CD
- size limit: 10GB images (Harbor)

Orchestration

- multi-container applications/services
- docker-compose
 - tool to automate command-line options for Docker
 - YAML style
 - mounting/binding of directories
 - networking
 - configuration settings
 - secrets
 - shell commands to be run when starting container
- Kubernetes
 - target: large clusters and services
 - typically: microservice architecture (each service a container)
 - automatic load-balancing, self-healing, scalability, ...

Wrap-up

- Questions?
- Feedback?
- time left to try out something yourselves?